



Postgres Plus Database Administration

Postgres Plus Advanced Server (PPAS) 9.5



Day 4 - Developer

Postgres Plus Advanced Server (PPAS) 9.5



Day 4-1

Procedural Languages

Objectives

- In this module you will learn:
 - PostgreSQL Procedural Languages
 - Introduction to PL/PGSQL
 - How it works
 - PL/pgSQL Block Structure
 - Declaring Variables
 - Writing Executable Statements
 - Declaring Function Parameters
 - Control Structures
 - Exception Handling
 - PL/pgSQL Cursors
 - Triggers
 - Examples & Lab

User-defined Functions

- PostgreSQL provides four kinds of functions:
 - query language functions (functions written in SQL)
 - procedural language functions (functions written in, for example, PL/pgSQL or PL/Tcl)
 - internal functions
 - C-language functions

```
CREATE [ OR REPLACE ] FUNCTION
    name ( [ [ argmode ] [ argname ] argtype [ { DEFAULT | = } default_expr ] [, ...] ] )
    [ RETURNS rettype
      | RETURNS TABLE ( column_name column_type [, ...] ) ]
    { LANGUAGE lang_name
      | TRANSFORM { FOR TYPE type_name } [, ... ]
      | WINDOW
      | IMMUTABLE | STABLE | VOLATILE | [ NOT ] LEAKPROOF
      | CALLED ON NULL INPUT | RETURNS NULL ON NULL INPUT | STRICT
      | [ EXTERNAL ] SECURITY INVOKER | [ EXTERNAL ] SECURITY DEFINER
      | COST execution_cost
      | ROWS result_rows
      | SET configuration_parameter { TO value | = value | FROM CURRENT }
      | AS 'definition'
      | AS 'obj_file', 'link_symbol'
    } ...
    [ WITH ( attribute [, ...] ) ]
```

Query Language Function

- SQL functions execute an arbitrary list of SQL statements.
- The first row of the last query's result will be returned.

```
CREATE FUNCTION tf1 (accountno integer, debit numeric) RETURNS integer AS $$  
    UPDATE bank  
        SET balance = balance - debit  
        WHERE accountno = tf1.accountno;  
    SELECT balance FROM bank WHERE accountno = tf1.accountno;  
$$ LANGUAGE SQL;  
  
SELECT tf1(17, 100.0);
```

- You can return multiple records using SETOF

```
CREATE FUNCTION getfoo(int) RETURNS SETOF foo AS $$  
    SELECT * FROM foo WHERE fooid = $1;  
$$ LANGUAGE SQL;  
  
SELECT * FROM getfoo(1) AS t1;
```

PostgreSQL Procedural Languages

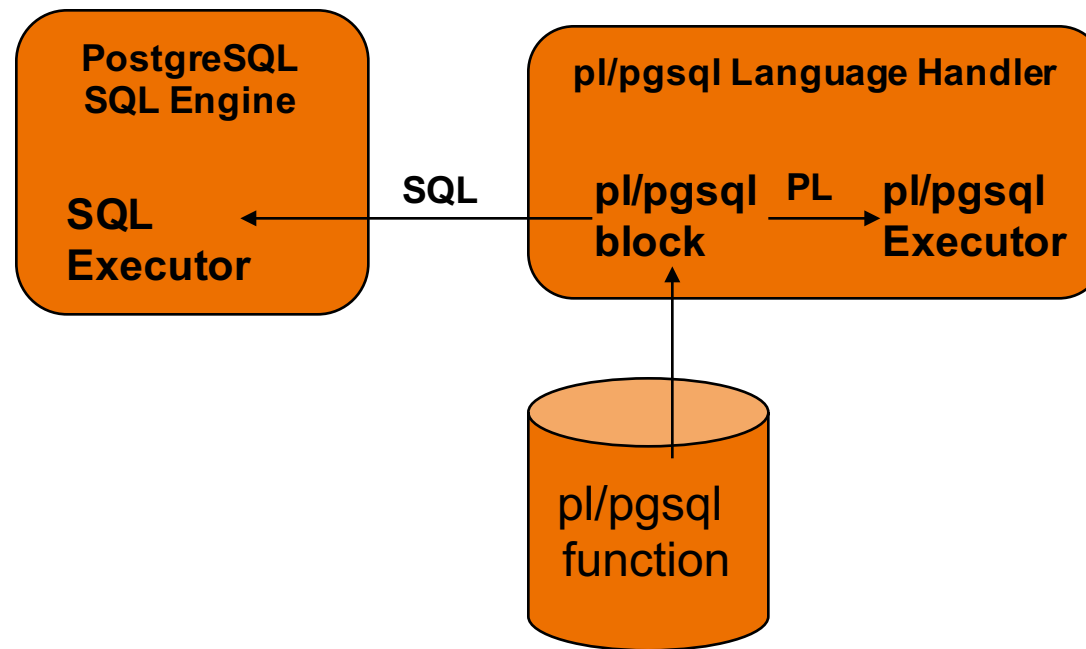
- PostgreSQL allows user-defined functions to be written in a variety of procedural languages.
- PostgreSQL currently supports several standard procedural languages :
 - PL/pgSQL
 - PL/Tcl
 - PL/Perl
 - PL/Python
 - PL/Java
 - PL/Ruby
 - Other languages can be defined by users

Introduction to PL/PGSQL

- PL/pgSQL is a loadable procedural language for the PostgreSQL database system.
- PL/pgSQL stand for Procedural language extension for postgres SQL
- PL/pgSQL has several distinct features:
 - Can be used to create functions and trigger procedures,
 - Adds control structures to the SQL language,
 - Can perform complex computations,
 - Inherits all user-defined types, functions, and operators,
 - Can be defined to be trusted by the server,
 - Is easy to use.

How it works

- PL/pgSQL defines a block structure for writing code
- PL/pgSQL functions can be compiled and stored inside database server offering better performance



PL/pgSQL Block Structure

- PL/pgSQL is a block-structured language.
- The complete text of a function definition must be a block.
 - DECLARE (optional)
 - Variables, cursors, user-defined exceptions
 - BEGIN (mandatory)
 - SQL statements
 - PL/pgSQL statements
 - EXCEPTION (optional)
 - Actions to perform when errors occur
 - END; (mandatory)
- Each declaration and each statement within a block is terminated by a semicolon. A block that appears within another block must have a semicolon after END, as shown above; however the final END that concludes a function body does not require a semicolon.

PL/pgSQL Structure

- All key words and identifiers can be written in mixed upper and lower case
- Identifiers are implicitly converted to lowercase unless double-quoted.
- There are two types of comments in PL/pgSQL.
 - -- starts a comment that extends to the end of the line.
 - /* multi-line comments */

Declaring Variables

Variables

- Temporary storage of data
- Can be reused in the plpgsql block
- Declared and initialized in the declarative section
- Used and assigned new values in the executable section
- Passed as parameters to PL/pgSQL subprograms
- Used to hold the output of a PL/pgSQL subprogram
- Can be any valid SQL data type
- Can contain a DEFAULT or CONSTANT clause

Declaring Variables

- Examples:
 - `user_id integer;`
 - `quantity numeric(5);`
 - `url varchar;`
 - `myrow tablename%ROWTYPE;`
 - `myfield tablename.columnname%TYPE;`
 - `arow RECORD;`
 - `quantity integer DEFAULT 32;`
- The general syntax of a variable declaration is:
 - `name [CONSTANT] type [NOT NULL] [{ DEFAULT | := } expression];`

Assign value to Variables

- Assignment

- `identifier := expression;`
- `user_id := 20;`
- `tax := subtotal * 0.06;`

- SELECT INTO

- `SELECT INTO target select_expressions FROM ...;`
- `SELECT INTO myrec * FROM emp WHERE empname = myname;`

Local Variable

FOUND

- Boolean local variable
- Starts out false within each PL/pgSQL function call.
- It is set by each of the following types of statements:
 - A SELECT INTO statement sets FOUND true if it returns a row, false if no row is returned
 - A PERFORM statement sets FOUND true if it produces (and discards) a row, false if no row is produced
 - UPDATE, INSERT, and DELETE statements set FOUND true if at least one row is affected, false if no row is affected
 - A FETCH statement sets FOUND true if it returns a row, false if no row is returned
 - A FOR statement sets FOUND true if it iterates one or more times, else false.

Writing Executable Statements

- Executable Statements are written inside BEGIN .. END block
- Blocks can be nested
- Statements inside a block can continue over several lines
- All statements must end with ;

Declaring Function Parameters

- Parameters passed to functions are named with the identifiers \$1, \$2, etc
- Optionally alias names can be used
- Either the alias or the numeric identifier can then be used
- E.g
 - `CREATE FUNCTION sales_tax(subtotal real)`

Example 1

- Following function will sum two numerics which are passed at function call:

```
create or replace function sum_val(i numeric, z numeric) returns  
numeric as $$
```

```
begin
```

```
return i+z;
```

```
end;
```

```
$$ language plpgsql;
```

- Invoke the function

```
Select sum_val(10,20);
```

Example 2

- Following function will show the difference between two passed values:

```
create or replace function differ(x numeric, y numeric) returns  
void as $$
```

```
declare
```

```
d numeric;
```

```
begin
```

```
D := x - y;
```

```
raise notice 'Difference is: %',d;
```

```
return;
```

```
end; $$ language plpgsql;
```

- Execute the function

```
select differ(10,5);
```

Anonymous code block - DO

DO [LANGUAGE *lang_name*] *code*

– *lang_name*: the default is plpgsql.

- execute an anonymous code block

```
DO $$DECLARE r record;
BEGIN
    FOR r IN SELECT table_schema, table_name FROM information_schema.tables
              WHERE table_type = 'VIEW' AND table_schema = 'public'
    LOOP
        EXECUTE 'GRANT ALL ON ' || quote_ident(r.table_schema) || '.' ||
quote_ident(r.table_name) || ' TO webuser';
    END LOOP;
END$$;
```

Control Structures

- Control Structures change the logical flow of statements with a plpgsql block
- Three main types of control structures are available in plpgsql:
 - IF statement
 - Case expressions
 - Loop control structures

IF statement

- plpgsql IF statement structure is similar to other procedural languages
- It allows plpgsql to perform actions selectively based on conditions

- **Syntax:**

```
IF boolean-expression THEN
statements
[ ELSIF boolean-expression THEN
statements]
[ ELSE
statements ]
END IF;
```

Example

- Following function will return account details from customer table for a given customerid:

```
create or replace function dis_cus(c_id numeric) returns void as $$
declare
rec RECORD;
begin
select into rec * from customer where custid=c_id;
if found then
raise notice 'Customer Name: %', rec.custname;
raise notice 'Account No:      %', rec.accountid;
else
raise notice 'No Data';
end if;
return;
end;$$ language plpgsql;
```

- Execute the function

```
Select dis_cus(130);
```

Case expression

- A CASE expression returns a result based on one or more alternatives
- The value of the selector determines which result is returned

- Syntax:

```
CASE search-expression
WHEN expression [, expression [ ... ]] THEN
statements
[ WHEN expression [, expression [ ... ]] THEN
statements
... ]
[ ELSE
statements ]
END CASE;
```


Example

- Following function will display temperature string based on input temperature reading using searched case:

```
CREATE OR REPLACE FUNCTION dis_temp(tem numeric) RETURNS varchar AS $$ DECLARE
msg varchar;
BEGIN
CASE WHEN tem < 0 THEN
msg := 'ICY';
WHEN tem between 0 and 10 THEN
msg := 'COLD';
WHEN tem > 10 THEN
msg := 'NORMAL';
ELSE
msg := 'Cannot determine';
END CASE;
return msg;
END; $$ language plpgsql;
```

- Execute the function:
select dis_temp(-10);

LOOP Statements

- Loops are mainly used to execute statements repeatedly until an exit condition is reached
- There are three loop types:
 - Basic loop
 - WHILE loop
 - FOR loop

Simple Loop

- Simple loop that performs repetitive actions without overall conditions
- A Simple loop must have an EXIT

- **Syntax:**

```
LOOP  
statements  
END LOOP;
```

- **Syntax for EXIT**

```
EXIT [ label ] [ WHEN boolean-expression ];
```

Example

- Following function will display a series 0-10 using simple loop:

```
create or replace function dis() returns void as $$
declare
rec numeric;
begin
rec:=0;
loop
raise notice '%',rec;
rec:=rec+1;
EXIT when rec>10;
end loop;
return;
end; $$ language plpgsql;
```

- Execute function:

```
Select dis();
```

While Loop

- While loops that perform iterative actions based on a condition
- The condition is evaluated at the start of each iteration
- If the condition yields NULL loop will exit

- Syntax:

```
WHILE boolean-expression LOOP  
statements  
END LOOP;
```

Example

- Following function will display a series 0-10 using while loop:

```
create or replace function dis() returns void as $$
declare
rec numeric;
begin
rec:=0;
while rec<=10 loop
raise notice '%',rec;
rec:=rec+1;
end loop;
return;
end; $$ language plpgsql;
```

- Execute function:

```
Select dis();
```

For Loop

- FOR loops that perform iterative actions based on a count
- FOR loops have the same general structure as the basic loop
- A control statement is used before the LOOP keyword to set the number of iterations
- Syntax:

```
FOR name IN [ REVERSE ] expression .. expression [ BY expression]  
  LOOP  
    statements  
  END LOOP;
```

Example

- Following function will display sequence of numbers 0-10 with increment of 2:

```
create or replace function dis1() returns void as
$$
declare
rec numeric;
begin
for rec in 0..10 by 2 loop
raise notice '%',rec;
end loop;
return;
end;$$ language plpgsql;
```

- Execute function:

```
Select dis1();
```


Example

- Following function will list name and account number for all the customers in customer table:

```
create or replace function dis_cus() returns void as
$$
declare
rec record;
begin
for rec in select*from customer loop
raise notice 'Customer Name: %', rec.custname;
raise notice 'Account No:      %', rec.accountid;
end loop;
return;
end;$$ language plpgsql;
```

- Execute function:
Select dis_cus();

Trapping Errors

- Syntax errors are handled at compile time
- plpgsql code may cause some unanticipated errors at run time
- To deal with such errors plpgsql provide EXCEPTION block

- **Syntax:**

```
[ DECLARE
declarations ]
BEGIN
statements
EXCEPTION
WHEN condition [ OR condition ... ] THEN
handler_statements
[ WHEN condition [ OR condition ... ] THEN
handler_statements
... ]
END;
```

Trapping Errors

- The EXCEPTION keyword starts the exception handling section
- Several exception handlers are allowed
- Only one handler is processed before leaving the block
- WHEN OTHERS is the last clause and can handle all types of exceptions
- EXCEPTION block must be at the last in a BEGIN..END block

Obtaining information about an error

- Exception handlers frequently need to identify the specific error that occurred.
- There are two ways to get information about the current exception:
 - Special variables
 - GET STACKED DIAGNOSTICS command.

Syntax: GET STACKED DIAGNOSTICS variable = item [, ...];

- Each item is a key word: RETURNED_SQLSTATE, MESSAGE_TEXT, PGPG_EXCEPTION_HINT, _EXCEPTION_DETAIL, PG_EXCEPTION_CONTEXT

Example

- Following function will divide first parameter with second and return the result:

```
create or replace function dis(a numeric, b numeric) returns numeric
as $$
```

```
declare
```

```
result numeric;
```

```
begin
```

```
result=a/b;
```

```
EXCEPTION
```

```
When others then
```

```
Raise notice 'Wrong value for second parameter. Must be a non-zero
value';
```

```
return result;
```

```
- end; $$ language plpgsql;
```

– **Execute function:**

```
Select dis();
```

PL/pgSQL Cursors

- Declaring Cursors

- `DECLARE curs1 refcursor;`
- `DECLARE curs2 CURSOR FOR SELECT * FROM tenk1;`
- `DECLARE curs3 CURSOR (key integer) IS SELECT * FROM tenk1 WHERE unique1 = key;`

- Opening Cursors

- `OPEN FOR query`
- `OPEN unbound_cursor FOR query;`
 - `OPEN curs1 FOR SELECT * FROM foo WHERE key = mykey;`
- `OPEN FOR EXECUTE`
- `OPEN unbound_cursor FOR EXECUTE query_string;`
 - `OPEN curs1 FOR EXECUTE 'SELECT * FROM ' || quote_ident($1);`

PL/pgSQL Cursors (cont)

- Opening a Bound Cursor

- `OPEN bound_cursor [(argument_values)] ;`
Opening Cursors
- `OPEN FOR query`
- DECLARE
 - `cur2 CURSOR FOR SELECT * FROM tenk1;`
 - `cur3 CURSOR (key integer) IS SELECT * FROM tenk1 WHERE
unique1 = key;`
- `OPEN cur2;`
- `OPEN cur3(42);`

PL/pgSQL Cursors (cont)

- FETCH
 - FETCH cursor INTO target;
 - FETCH retrieves the next row from the cursor into a target, which may be a row variable, a record variable, or a comma-separated list of simple variables, just like SELECT INTO.
 - As with SELECT INTO, the special variable FOUND may be checked to see whether a row was obtained or not.
 - `FETCH curs1 INTO rowvar;`
 - `FETCH curs2 INTO foo, bar, baz;`
- CLOSE
 - CLOSE closes the portal underlying an open cursor.
 - This can be used to release resources earlier than end of transaction, or to free up the cursor variable to be opened again.
 - `CLOSE curs1;`

Example

- Following function will list name and account number for all the customers in customer table:

```
- create or replace function dis_cus() returns void as $$  
- declare  
- cur CURSOR for select*from customer;  
- rec RECORD;  
- begin  
- open cur;  
- fetch cur into rec;  
- while found loop  
- raise notice 'Customer Name: %', rec.custname;  
- raise notice 'Account No:      %', rec.accountid;  
- fetch cur into rec;  
- end loop;  
- close cur;  
- return;  
- end; $$ language plpgsql;
```

- Execute function:

```
- Select dis_cus();
```

Triggers

- Created with the CREATE FUNCTION command
- Several special variables are created automatically in the top-level block.
 - NEW
 - Data type RECORD; variable holding the new database row for INSERT/UPDATE operations in row-level triggers. This variable is NULL in statement-level triggers.
 - OLD
 - Data type RECORD; variable holding the old database row for UPDATE/DELETE operations in row-level triggers. This variable is NULL in statement-level triggers.
 - TG_NAME
 - Data type name; variable that contains the name of the trigger actually fired.

Triggers

- TG_WHEN
 - Data type text; a string of either BEFORE or AFTER depending on the trigger's definition.
- TG_LEVEL
 - Data type text; a string of either ROW or STATEMENT depending on the trigger's definition.
- TG_OP
 - Data type text; a string of INSERT, UPDATE, or DELETE telling for which operation the trigger was fired.
- TG_RELNAME
 - Data type name; the name of the table that caused the trigger invocation.
- TG_NARGS
 - Data type integer; the number of arguments given to the trigger procedure in the CREATE TRIGGER statement.

Triggers

- Note: A trigger function must return either NULL or a record/row value having exactly the structure of the table the trigger was fired for.
- The return value of a BEFORE or AFTER statement-level trigger or an AFTER row-level trigger is always ignored; it may as well be null. However, any of these types of triggers can still abort the entire operation by raising an error.

Example

- In this example city table store total population of a city and population will increase by 1 when a birth registration is done for that city:
- Create base tables:
 - `create table city(cityid numeric, population numeric);`
 - `create table birth(regid numeric, name varchar, cityid numeric);`
 - `insert into city values(1,0),(2,0);`
 - `select*from city;`
- Create Trigger function that will be called by trigger:
create or replace function trg_ins() returns trigger as \$\$
begin
update city set population=population+1 where cityid=NEW.cityid;
return null;
end; \$\$ language plpgsql;
- Create Trigger on birth table:
create trigger trg_ins_brth after insert on birth
for each row execute procedure trg_ins();
- Test whether trigger is working:
`insert into birth values(101,'Raj',2);`
`select*from birth;`
`select*from city;`

Lab Exercise - 1

- Create a plpgsql function
 - Declare two variables in declare section
 - Assign today's date to the first variable
 - Assign tomorrow's date to the second variable
 - Return first variable to the calling context
 - Display tomorrow's date on screen using notice messages

Lab Exercise - 2

- DECLARE
 citizen VARCHAR(50) := 'SENIOR';
 credit_rating VARCHAR(50) := 'EXCELLENT';
BEGIN
 DECLARE
 citizen NUMBER(7) := 201;
 name VARCHAR(25) := 'JUNIOR';
 BEGIN
 credit_rating := 'GOOD';
 ...
 END;
 ...
END;
- In the plpgsql function above determine the value of following variable:
 - citizen in nested block:
 - credit_rating in main block:
 - citizen in main block:
 - name in main block:

Lab Exercise - 3

- Create a plpgsql function which can be called with a customerid and display hello and firstname of that customer
- Connect with edbstore database and create a table using following statement:
- create table saleproducts (like products);
- Create a plpgsql function to copy all products priced 9.99,10.99,11.99 from products table to saleproducts table

Lab Exercise - 4

- Create a plpgsql function which can be called with price as input and delete all the rows from saleproducts table for given price tag
- Create a plpgsql function using cursors which display category and highest priced prod_id in that category

PL/Python

- Add environment variables

```
Export PYTHONHOME=/opt/EnterpriseDB/LanguagePack/9.5/Python-3.3

export PATH=$PYTHONHOME/bin:$PATH
export LD_LIBRARY_PATH=$PYTHONHOME/lib:$LD_LIBRARY_PATH
```

- How to install

```
postgres@pglab1:~$ psql -c 'CREATE LANGUAGE plpython3u;'
CREATE LANGUAGE
```

- Example

```
CREATE FUNCTION gethostbyname(hostname text)
    RETURNS inet
AS $$
    import socket
    return socket.gethostbyname(hostname)
$$ LANGUAGE plpython3u SECURITY DEFINER;
```

- Result

```
postgres=# SELECT gethostbyname('www.postgresql.org');
 gethostbyname
-----
 87.238.57.232
(1 row)
```

PL/Python Table functions

```
CREATE FUNCTION even_numbers_from_list(up_to int)
  RETURNS SETOF int
AS $$
    return range(0,up_to,2)
$$ LANGUAGE plpython3u;
```

```
CREATE OR REPLACE FUNCTION even_numbers_from_generator(up_to int)
  RETURNS TABLE (even int, odd int)
AS $$
    return ((i,i+1) for i in range(0,up_to,2))
$$ LANGUAGE plpython3u;
```

```
CREATE OR REPLACE FUNCTION even_numbers_with_yield(up_to int,
                                                    OUT even int, OUT odd int)
  RETURNS SETOF RECORD
AS $$
    for i in range(0,up_to,2):
        yield i, i+1
$$ LANGUAGE plpython3u;
```

PL/Python Query

- `plpy.execute`, `plpy.prepare`

```
CREATE OR REPLACE FUNCTION get_table_oid(table_name text)
    RETURNS OID
AS $$
    stmt = plpy.prepare("select oid from pg_class where relname = $1", ["text"])
    res = plpy.execute(stmt, [table_name])
    return res[0]["oid"]
$$ LANGUAGE plpython3u;
```

```
postgres=# select get_table_oid('test');
 get_table_oid
-----
          16408
(1 row)
```

PL/Python Logging

- 화면 출력을 위해서는 `plpy.notice()` 함수를 사용

```
CREATE OR REPLACE FUNCTION fact(x int) RETURNS int
AS $$
    global x
    f = 1
    while (x > 0):
        f = f * x
        x = x - 1
        plpy.notice('f:%d, x:%d' % (f, x))
    return f
$$ LANGUAGE plpython3u;
```

```
postgres=# select * from fact(3);
NOTICE:  f:3, x:2
CONTEXT:  PL/Python function "fact"
NOTICE:  f:6, x:1
CONTEXT:  PL/Python function "fact"
NOTICE:  f:6, x:0
CONTEXT:  PL/Python function "fact"
fact
-----
      6
(1 row)
```

Summary

- In this module you learned:
 - PostgreSQL Procedural Languages
 - Introduction to PL/PGSQL
 - How it works
 - PL/pgSQL Block Structure
 - Declaring Variables
 - Writing Executable Statements
 - Declaring Function Parameters
 - Control Structures
 - Exception Handling
 - PL/pgSQL Cursors
 - Triggers
 - Examples & Lab



Day 4-2

Extension Development

Writing Functions in C - Example

- add_func.c

```
#include "postgres.h"
#include "fmgr.h"

PG_MODULE_MAGIC;

PG_FUNCTION_INFO_V1(add_ab);

Datum
add_ab(PG_FUNCTION_ARGS)
{
    int32    arg_a = PG_GETARG_INT32(0);
    int32    arg_b = PG_GETARG_INT32(1);

    PG_RETURN_INT32(arg_a + arg_b);
}
```

- Makefile

```
MODULES = add_func

PG_CONFIG = pg_config
PGXS := $(shell $(PG_CONFIG) --pgxs)
include $(PGXS)
```

- Compile

```
pglab@pglab1:~$ sudo su -
root@pglab1:~# cd ~postgres/temp
root@pglab1:/opt/PostgreSQL/9.4/temp# ls
add_func.c  Makefile
root@pglab1:/opt/PostgreSQL/9.4/temp# make
gcc -Wall ... -o add_func.o add_func.c
gcc -Wall ... -o add_func.so add_func.o
root@pglab1:/opt/PostgreSQL/9.4/temp# ls
add_func.c  add_func.o  add_func.so  Makefile
root@pglab1:/opt/PostgreSQL/9.4/temp# make install
/bin/mkdir -p '/opt/PostgreSQL/9.4/lib/postgresql'
/usr/bin/install -c -m 755 add_func.so
'/opt/PostgreSQL/9.4/lib/postgresql/'
```


Create function using C library

```
CREATE OR REPLACE FUNCTION add(int, int) RETURNS INT
AS '/opt/PostgreSQL/9.4/lib/postgresql/add_func', 'add_ab'
LANGUAGE C STRICT;
```

```
postgres=# select add(1, 20);
 add
-----
  21
(1 row)
```

```
postgres=# \df+ add
```

List of functions										
Schema	Name	Result data type	Argument data types	Type	Security	Volatility	Owner	Language	Source code	Description
public	add	integer	integer, integer	normal	invoker	volatile	postgres	c	add_ab	

(1 row)

add_funcs.sql.in

- CREATE FUNCTION 구문 자동 생성

```
CREATE OR REPLACE FUNCTION add(int, int) RETURNS INT
AS 'MODULE_PATHNAME', 'add_ab'
LANGUAGE C STRICT;
```

- Makefile 에 아래와 같이 추가

```
MODULES = add_func
DATA_built = add_funcs.sql
```

```
PG_CONFIG = pg_config
PGXS := $(shell $(PG_CONFIG) --pgxs)
include $(PGXS)
```

- Build

```
root@pglab1:/opt/PostgreSQL/9.4/temp# ls
add_func.c  add_funcs.sql.in  Makefile
root@pglab1:/opt/PostgreSQL/9.4/temp# make
sed 's,MODULE_PATHNAME,$libdir/add_funcs,g' add_funcs.sql.in >add_funcs.sql
...
root@pglab1:/opt/PostgreSQL/9.4/temp# cat add_funcs.sql
CREATE OR REPLACE FUNCTION add(int, int) RETURNS INT
AS '$libdir/add_funcs', 'add_ab'
LANGUAGE C STRICT;
```

Custom Type

- PostgreSQL의 Type system은 쉽게 확장이 가능
- 정해진 몇개의 C function을 구현하면 누구나 구현 가능
- 타 DB와 차별화 되는 PostgreSQL 만의 특 장점
- Complex number 예제

```
postgres=# \d test_complex
Table "public.test_complex"
Column | Type      | Modifiers
-----+-----+-----
a      | complex   |
b      | complex   |
Indexes:
    "test_cplx_ind" btree (a)
```

```
postgres=# select * from test_complex;
      a      |      b
-----+-----
(1,2.5)      | (4.2,3.55)
(33,51.4)     | (100.42,93.55)
(56,-22.5)    | (-43.2,-0.07)
(-91.9,33.6) | (8.6,3)
(4 rows)
```

```
postgres=# select a+b abs_sum from test_complex;
      abs_sum
-----
(5.2,6.05)
(133.42,144.95)
(12.8,-22.57)
(-83.3,36.6)
(4 rows)
```

Sample – 주민번호 Type

- 주민번호를 저장하기 위한 jumin 타입
- 기능
 - 주민번호 뒷자리 암호화 및 마스킹
 - Equality 비교
 - 생년월일 순 정렬 및 비교
 - 주민번호 뒷자리 중 첫 글자를 기준으로 1800, 1900, 2000년대 자동 처리 기능
 - 생년월일 추출
 - 성별 추출

주민번호 Type

```
CREATE TABLE test_jumin (  
    a    jumin,  
    b    char(14)  
);
```

```
INSERT INTO test_jumin VALUES ('951010-9723134', '951010-9723134');  
INSERT INTO test_jumin VALUES ('750107-1253452', '750107-1253452');  
INSERT INTO test_jumin VALUES ('750107-1253452', '750107-1253452');  
INSERT INTO test_jumin VALUES ('900208-2575294', '900208-2575294');
```

```
postgres=# select * from test_jumin order by a;
```

a	b
951010-9*****	951010-9723134
951010-9*****	951010-9723134
951010-9*****	951010-9723134
971101-0*****	971101-0723134
731218-1*****	731218-1151958
750107-1*****	750107-1253452
750107-1*****	750107-1253452
760907-1*****	760907-1756006
761209-2*****	761209-2436370
830207-2*****	830207-2187272
870520-2*****	870520-2508631
880205-1*****	880205-1163291
900208-2*****	900208-2575294
951227-2*****	951227-2438973
960806-2*****	960806-2405962
960806-2*****	960806-2205963

(16 rows)

```
postgres=# select * from test_jumin order by b;
```

a	b
731218-1*****	731218-1151958
750107-1*****	750107-1253452
750107-1*****	750107-1253452
760907-1*****	760907-1756006
761209-2*****	761209-2436370
830207-2*****	830207-2187272
870520-2*****	870520-2508631
880205-1*****	880205-1163291
900208-2*****	900208-2575294
951010-9*****	951010-9723134
951010-9*****	951010-9723134
951010-9*****	951010-9723134
951227-2*****	951227-2438973
960806-2*****	960806-2205963
960806-2*****	960806-2405962
971101-0*****	971101-0723134

(16 rows)

주민번호 Type - Cont

```
postgres=# select * from test_jumin where b like '960806-2%';
```

a	b
960806-2*****	960806-2405962
960806-2*****	960806-2205963

(2 rows)

```
postgres=# select * from test_jumin where a = '960806-2405962';
```

a	b
960806-2*****	960806-2405962

(1 row)

```
postgres=# select a, jumin_birthday(a), jumin_gender(a)
```

a	jumin_birthday	jumin_gender
951010-9*****	18951010	M
951010-9*****	18951010	M
951010-9*****	18951010	M
971101-0*****	18971101	F
731218-1*****	19731218	M
750107-1*****	19750107	M
750107-1*****	19750107	M
760907-1*****	19760907	M
761209-2*****	19761209	F
830207-2*****	19830207	F
870520-2*****	19870520	F
880205-1*****	19880205	M
900208-2*****	19900208	F
951227-2*****	19951227	F
960806-2*****	19960806	F
960806-2*****	19960806	F

(16 rows)

What are Extension modules?

- Modules in the “Extension” directory are additional features to PostgreSQL.
- Generally, Extension modules aren’t included in the core database because they don’t appeal to a wide-enough audience, or because they are still under development.
- Extension modules are sometimes pulled into the core distribution.
- The modules in Extension are tied to that particular version of PostgreSQL, and the modules that are available change over time.

Installing Extension Modules

- In order to use a Extension module in a database you need to run `CREATE EXTENSION` command to install the module's features into that database.
- By default, Extension files are installed at `PREFIX/share/extension`, but most package management systems change this to `share/PostgreSQL/extension` or some variant.
- Documentation for Extension modules is installed at `share/Extension` or equivalent.

Installing Extension Modules

- You need to register the new objects in the database system by running create extension command
- Alternatively, run it in database template1 so that the module will be copied into subsequently-created databases by default.
- Syntax:

```
CREATE EXTENSION module_name;
```

-This command must be run by a database superuser



Day 4-3

Connectors

Objectives

- This module will cover:
 - Installing JDBC Connectors
 - Installing .NET Connectors
 - Extensions
 - Other resources

PostgreSQL JDBC driver

- JDBC API defines how clients may access the database
- Postgres JDBC Drivers provide connection to Postgres for Java applications
- Drivers available in for of Jar files
- Download from <http://jdbc.postgresql.org/download.html>
- Use JDBC4 for JDK1.6 or higher
- Other versions of JDBC drivers also available
- Source code also available

Connecting Java to Postgres

- Step 1: Import java.sql
- Step 2: Load the drivers `Class.forName("org.postgresql.Driver");`
- Step 3: Prepare Connection URL
 - jdbc:postgresql:database
 - jdbc:postgresql://host/database
 - jdbc:postgresql://host:port/database
- Step 4: get a Connection instance from JDBC
 - `Connection con = DriverManager.getConnection(url, username, password);`

EDB Postgres JDBC driver

- The EDB JDBC driver is a super set of community JDBC driver.
- Provides additional features for Oracle compatibilities.
 - It's highly recommended to use EDB driver with Oracle compatible mode.
- <http://www.enterprisedb.com/docs/en/9.5/jdbc/toc.html>
- Drivers available in for of Jar files
 - edb-jdbc15.jar supports JDBC version 3
 - edb-jdbc16.jar supports JDBC version 4
 - edb-jdbc17.jar supports JDBC version 4.1
- Driver class name is different
 - com.edb.Driver (not com.postgresql.Driver)
- Conn string scheme is different
 - jdbc:edb:database
 - jdbc:edb://host/database
 - jdbc:edb://host:port/database

Connection Parameters

- Example of Various Connection Properties:
 - `String url = "jdbc:postgresql://localhost/edbstore";`
 - `Properties prt = new Properties();`
 - `prt.setProperty("user","postgres");`
 - `prt.setProperty("password","postgres1");`
 - `prt.setProperty("ssl","true");`
 - `Connection conn = DriverManager.getConnection(url, prt);`

 - `String url =`
`"jdbc:postgresql://localhost/edbstore?user=postgres&password=postgres1&ssl=true";`
 - `Connection conn = DriverManager.getConnection(url);`

JDBC CopyManager

- Normal insert
 - 800 seconds
- JDBC Batch update
 - 25 seconds
- Copy Manager
 - 7 seconds

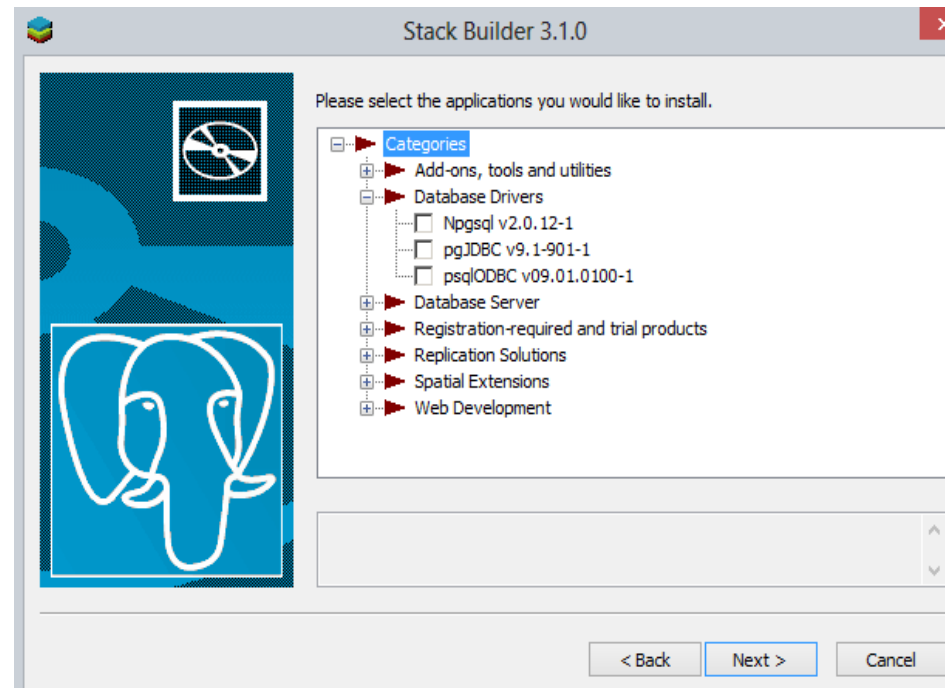
```
CopyManager cm = new CopyManager((BaseConnection) conn);
cpIN= cm.copyIn("COPY test(id, dat1, dat2) FROM STDIN WITH DELIMITER '|'");
StringBuffer buf = new StringBuffer();
byte[] data = buf.append(i).append("|ABCDE|ABCDEFGHijklmn\n").toString().getBytes();
cpIN.writeToCopy(data, 0, data.length);
cpIN.endCopy();
pyManager cm = new CopyManager((BaseConnection) conn);
```


.NET Connectors

- Microsoft .NET CLR can connect to Postgres using:
 - Npgsql .NET Data Provider
 - Via ODBC Data Source
- Allows .NET application to connect with Postgres database
- Easy installation using Stack Builder from EnterpriseDB
- Connector Sources are available
- <http://www.enterprisedb.com/docs/en/9.5/dotnet/toc.html>

.NET Connectors Installation

- Open Stack Builder from start menu and click next after select your Postgres installation
- Expand database drivers node and select the npgsql or psqLODBC



.NET Npgsql Driver Features

- You can send select, insert, delete queries
- You can call functions
- You can get resultset from functions
- You can use parameters in your queries. Input, Output and InputOutput parameters are supported
- Support for transactions

Connecting .NET to Postgres

- Npgsql data access is very similar to .NET connecting to SQL Server or OleDb
 - Add required namespace: Using Npgsql
 - Gather Connection String parameters information
 - Establish Connection:
 - `NpgsqlConnection conn = new NpgsqlConnection("Server=127.0.0.1;Port=5432;UserId=postgres;Password=postgres1;Database=edbstore;");`
 - Open connect using `open()`
 - Use `NpgsqlCommand` class to create a query:
 - `NpgsqlCommand command = new NpgsqlCommand("insert into sample values(1, 1)", conn)`
 - Execute Query using `ExecuteNonQuery()`
 - Close connection using `Close()`



Day 4-4

Useful Data Types

bytea

- The bytea data type allows storage of binary strings. BLOB

```
edb=# select 'abcde', 'abcde'::bytea;
?column? |      bytea
-----+-----
abcde    | \x6162636465
```

- Easily convert between types.

```
edb=# select '안녕하세요'::bytea;
      bytea
-----
\xec9588eb8595ed9598ec84b8ec9a94

edb=# select convert('안녕하세요', 'utf8', 'euc-kr');
      convert
-----
\xbec8b3e7c7cfbcbcbfe4

edb=# select convert_from(E'\xbec8b3e7c7cfbcbcbfe4'::bytea, 'euc-kr');
      convert_from
-----
안녕하세요
```

Lab

Network Address Types

Name	Storage Size	Description
cidr	7 or 19 bytes	IPv4 and IPv6 networks
inet	7 or 19 bytes	IPv4 and IPv6 hosts and networks
macaddr	6 bytes	MAC addresses

- Data types to store IPv4, IPv6, and MAC addresses
- cidr is more strict. Validate subnet.

```
edb=# select cidr '192.168.56';
      cidr
-----
192.168.56.0/24

edb=# select inet '192.168.56.200/24';  # invalid subnet
      inet
-----
192.168.56.200/24

edb=# select cidr '192.168.56.200/24';
ERROR:  invalid cidr value: "192.168.56.200/24"
LINE 1: select cidr '192.168.56.200/24';
                        ^
DETAIL:  Value has bits set to right of mask.
```


cidr and inet Containment test

- Containment test

Operator	Description	Example
<<	is contained by	inet '192.168.1.5' << inet '192.168.1/24'
<=<	is contained by or equals	inet '192.168.1/24' <=< inet '192.168.1/24'
>>	contains	inet '192.168.1/24' >> inet '192.168.1.5'
>=>	contains or equals	inet '192.168.1/24' >=> inet '192.168.1/24'
&&	contains or is contained by	inet '192.168.1/24' && inet '192.168.1.80/28'

```
edb=# select cidr '192.168.56.192/26' >> inet '192.168.56.200';
?column?
-----
t
(1 row)


edb=# select cidr '192.168.56.192/26' >> inet '192.168.56.100';
?column?
-----
f
(1 row)
```

cidr and inet functions

Function	Return Type	Description	Example	Result
<code>abbrev(inet)</code>	text	abbreviated display format as text	<code>abbrev(inet '10.1.0.0/16')</code>	10.1.0.0/16
<code>abbrev(cidr)</code>	text	abbreviated display format as text	<code>abbrev(cidr '10.1.0.0/16')</code>	10.1/16
<code>broadcast(inet)</code>	inet	broadcast address for network	<code>broadcast('192.168.1.5/24')</code>	192.168.1.255/24
<code>family(inet)</code>	int	extract family of address; 4 for IPv4, 6 for IPv6	<code>family('::1')</code>	6
<code>host(inet)</code>	text	extract IP address as text	<code>host('192.168.1.5/24')</code>	192.168.1.5
<code>hostmask(inet)</code>	inet	construct hostmask for network	<code>hostmask('192.168.23.20/30')</code>	0.0.0.3
<code>masklen(inet)</code>	int	extract netmask length	<code>masklen('192.168.1.5/24')</code>	24
<code>netmask(inet)</code>	inet	construct netmask for network	<code>netmask('192.168.1.5/24')</code>	255.255.255.0
<code>network(inet)</code>	cidr	extract network part of address	<code>network('192.168.1.5/24')</code>	192.168.1.0/24
<code>set_masklen(inet, int)</code>	inet	set netmask length for inet value	<code>set_masklen('192.168.1.5/24', 16)</code>	192.168.1.5/16
<code>set_masklen(cidr, int)</code>	cidr	set netmask length for cidr value	<code>set_masklen('192.168.1.0/24'::cidr, 16)</code>	192.168.0.0/16
<code>text(inet)</code>	text	extract IP address and netmask length as text	<code>text(inet '192.168.1.5')</code>	192.168.1.5/32
<code>inet_same_family(inet, inet)</code>	boolean	are the addresses from the same family?	<code>inet_same_family('192.168.1.5/24', '::1')</code>	FALSE
<code>inet_merge(inet, inet)</code>	cidr	the smallest network which includes both of the given networks	<code>inet_merge('192.168.1.5/24', '192.168.2.5/24')</code>	192.168.0.0/22

macaddr

- The macaddr type stores MAC addresses. 6bytes
- Supported input formats



```
edb=# select macaddr '08002b:010203', macaddr '08-00-2b-01-02-03';
      macaddr      |      macaddr
-----+-----
 08:00:2b:01:02:03 | 08:00:2b:01:02:03
(1 row)
edb=# select trunc(macaddr '08002b:010203');
      trunc
-----
 08:00:2b:00:00:00
```

Lab

UUID

- Stores Universally Unique Identifiers (UUID). 16bytes.
- Supported input formats

```
A0EEBC99-9C0B-4EF8-BB6D-6BB9BD380A11  
{a0eebc99-9c0b-4ef8-bb6d-6bb9bd380a11}  
a0eebc999c0b4ef8bb6d6bb9bd380a11  
a0ee-bc99-9c0b-4ef8-bb6d-6bb9-bd38-0a11  
{a0eebc99-9c0b4ef8-bb6d6bb9-bd380a11}
```



```
a0eebc99-9c0b-4ef8-bb6d-6bb9bd380a11
```

```
edb=# select uuid 'A0EEBC99-9C0B-4EF8-BB6D-6BB9BD380A11';  
          uuid  
-----  
a0eebc99-9c0b-4ef8-bb6d-6bb9bd380a11  
  
edb=# create extension pgcrypto;  
CREATE EXTENSION  
  
edb=# select gen_random_uuid();  
          gen_random_uuid  
-----  
58deed14-e3b8-4799-b394-6e182eb66626
```

Lab



Day 4-5
JSON & JSONB

Why NoSQL?

- Where did NoSQL come from?
 - Where all cool tech stuff comes from – Internet companies
- Why did they make NoSQL?
 - To support huge data volumes and evolving demands for ways to work with new data types
- What does NoSQL accomplish?
 - Enables you to work with new data types: email, mobile interactions, machine data, social connections
 - Enables you to work in new ways: incremental development and continuous release
- Why did they have to build something new?
 - There were limitations to most relational databases

Postgres' Response

- HSTORE
 - Key-value pair
 - Simple, fast and easy
 - Postgres v 8.2 – pre-dates many NoSQL-only solutions
 - Ideal for flat data structures that are sparsely populated
- JSON
 - Hierarchical document model
 - Introduced in Postgres 9.2, perfected in 9.3
- JSONB
 - Binary version of JSON
 - Faster, more operators and even more robust
 - Postgres 9.4

HStore: Key-value Store

- Supported since 2006, the HStore contrib module enables storing key/value pairs within a single column
- Allows you to create a schema-less, ACID compliant data store within Postgres
- Create single HStore column and include, for each row, only those keys which pertain to the record
- Add attributes to a table and query without advance planning
- Combines flexibility with ACID compliance

HSTORE Examples

- Create a table with HSTORE field

```
CREATE TABLE hstore_data (data HSTORE);
```

- Insert a record into hstore_data

```
INSERT INTO hstore_data (data) VALUES ( '
    "cost"=>"500",
    "product"=>"iphone",
    "provider"=>"apple" ' );
```

- Select data from hstore_data

```
SELECT data FROM hstore_data ;
```

```
-----
"cost"=>"500", "product"=>"iphone", "provider"=>"Apple "
```

```
(1 row)
```

```
SELECT data->'product' FROM hstore_data;
```

```
-----
```

```
iphone
```

```
(1 row)
```

Postgres: Document Store

- JSON is the most popular data-interchange format on the web
- Derived from the ECMAScript Programming Language Standard.
- Supported by virtually every programming language
- New supporting technologies continue to expand JSON's utility
 - PL/V8 JavaScript extension
 - Node.js
- Postgres has a native JSON data type (v9.2) and a JSON parser and a variety of JSON functions (v9.3)
- Postgres will have a JSONB data type with binary storage and indexing (coming – v9.4)

JSON Examples

- Creating a table with a JSONB field

```
CREATE TABLE json_data (data JSONB);
```

- Simple JSON data element:

```
{"name": "Apple Phone", "type": "phone", "brand": "ACME", "price": 200,  
"available": true, "warranty_years": 1}
```

- Inserting this data element into the table json_data

```
INSERT INTO json_data (data) VALUES  
(' { "name": "Apple Phone",  
      "type": "phone",  
      "brand": "ACME",  
      "price": 200,  
      "available": true,  
      "warranty_years": 1  
    } ')
```

JSON Examples

- JSON data element with nesting:

```
{ "full name": "John Joseph Carl Salinger",  
  "names":  
    [  
      { "type": "firstname", "value": "John" },  
      { "type": "middlename", "value": "Joseph" },  
      { "type": "middlename", "value": "Carl" },  
      { "type": "lastname", "value": "Salinger" }  
    ]  
}
```

A simple query for JSON data

```
SELECT DISTINCT
    data->>'name' as products
FROM json_data;
```

products

```
-----
Cable TV Basic Service Package
AC3 Case Black
Phone Service Basic Plan
AC3 Phone
AC3 Case Green
Phone Service Family Plan
AC3 Case Red
AC7 Phone
```

This query does not return JSON data – it returns text values associated with the key 'name'

A query that returns JSON data

```
SELECT data FROM json_data;  
data
```

This query returns
the JSON data in its
original format

```
-----  
{ "name": "Apple Phone", "type":  
  "phone", "brand": "ACME", "price":  
  200, "available": true,  
  "warranty_years": 1 }
```


JSON and ANSI SQL - PB&J for the DBA

- JSON is naturally integrated with ANSI SQL in Postgres
- JSON and SQL queries use the same language, the same planner, and the same ACID compliant transaction framework
- JSON and HSTORE are elegant and easy to use extensions of the underlying object-relational model

JSON and ANSI SQL Example

```
SELECT DISTINCT
    product_type,
    data->>'brand' as Brand,
    data->>'available' as Availability
FROM json_data
JOIN products
ON (products.product_type=json_data.data->>'name')
WHERE json_data.data->>'available'=true;
```

product_type	brand	availability
AC3 Phone	ACME	true

No need for programmatic logic to combine SQL and NoSQL in the application

Bridging between SQL and JSON

Simple ANSI SQL Table Definition

```
CREATE TABLE products (id integer, product_name text );
```

Select query returning standard data set

```
SELECT * FROM products;
```

id	product_name
1	iPhone
2	Samsung
3	Nokia

Select query returning the same result as a JSON data set

```
SELECT ROW_TO_JSON(products) FROM products;
```

```
{"id":1,"product_name":"iPhone"}  
{"id":2,"product_name":"Samsung"}  
{"id":3,"product_name":"Nokia"}
```

JSON Data Type Example

```
{
  "firstName": "John",           -- String Type
  "lastName": "Smith",          -- String Type
  "isAlive": true,               -- Boolean Type
  "age": 25,                     -- Number Type
  "height_cm": 167.6,           -- Number Type
  "address": {                  -- Object Type
    "streetAddress": "21 2nd Street",
    "city": "New York",
    "state": "NY",
    "postalCode": "10021-3100"
  },
  "phoneNumbers": [             // Object Array
    {                             // Object
      "type": "home",
      "number": "212 555-1234"
    },
    {
      "type": "office",
      "number": "646 555-4567"
    }
  ],
  "children": [],
  "spouse": null                 // Null
}
```

JSONB

- Canonical representation
 - Whitespace and punctuation dissolved away
 - Only one value per object key is kept
 - Last insert wins
 - Key order determined by length, then bitwise comparison
- Equality, containment and key/element presence tests
- New JSONB creation functions
- Smaller, faster GIN indexes
- jsonb subdocument indexes
 - Use “get” operators to construct expression indexes on subdocument:
 - `CREATE INDEX author_index ON books USING GIN ((jsondata -> 'authors'));`
 - `SELECT * FROM books WHERE jsondata -> 'authors' ? 'Carl Bernstein'`

JSON, JSONB or HSTORE?

- JSON/JSONB is more versatile than HSTORE
- HSTORE provides more structure
- JSON or JSONB?
 - if you need any of the following, use JSON
 - Storage of validated json, without processing or indexing it
 - Preservation of white space in json text
 - Preservation of object key order Preservation of duplicate object keys
 - Maximum input/output speed
- For any other case, use JSONB

Lab



Day 4-6
Foreign Data Wrapper

Foreign Data Wrapper

- Foreign data wrappers (SQL MED) allow queries to read and write data to foreign data sources.
- Foreign data wrappers allow SQL access to data in:
 - Postgres clusters on the same server, perhaps in different databases or clusters
 - Remote Postgres servers
 - Stored in non-Postgres data repositories
 - Stored in data repositories with different performance and storage characteristics

Why use Foreign Data Wrapper?

- Using PostgreSQL as a central interface to connect to other systems / databases to gather data and issue queries or joins.
- Push-down for **WHERE** and **Column** to improve performance
- SELECT syntax; including useful clauses like DISTINCT, ORDER BY, GROUP BY and more.
- JOIN external data with internal tables
- FUNCTIONS for comparison, math, string, pattern matching, date/time, etc
- DML Operations INSERT / UPDATE / DELETE
- Make external data sources look like local tables

*The detail of all the available foreign data wrapper available
https://wiki.postgresql.org/wiki/Foreign_data_wrappers*

How to configure FDW?

- Creating an extension.

```
CREATE EXTENSION dummy_fdw;
```

- Create a foreign server.

```
CREATE SERVER dummy_server  
FOREIGN DATA WRAPPER dummy_fdw  
OPTIONS (host '127.0.0.1', port '3306');
```

- Creating a User-Mappings

```
CREATE USER MAPPING FOR postgres  
SERVER dummy_server  
OPTIONS (username 'foo', password 'bar');
```

- Creating a foreign table

```
CREATE FOREIGN TABLE f_table(id int, name text) SERVER dummy_server  
OPTIONS (dbname 'db', table_name 'r_table');
```

How to Use FDW?

- The usage of Foreign Table is almost the same as the PostgreSQL table.

- **Selecting Data From foreign table**

```
SELECT id, name FROM f_table WHERE id = 1;
```

- **Inserting data into foreign table.**

```
INSERT INTO f_table values (1, 'foor',);
```

```
INSERT INTO f_table values (2, 'bar',);
```

- **Deleting data from Foreign Table**

```
DELETE FROM f_table where id= 3;
```

- **Update Foreign Table**

```
UPDATE f_table set name= 'bar' WHERE id = 1;
```

- **Explain Table**

```
EXPLAIN SELECT id, name FROM f_table WHERE name LIKE 'foo' limit 1;
```

FDW internal

```
extern Datum ifx_fdw_handler(PG_FUNCTION_ARGS);
extern Datum ifx_fdw_validator(PG_FUNCTION_ARGS);

CREATE FUNCTION ifx_fdw_handler() RETURNS fdw_handler
  AS 'MODULE_PATHNAME'
  LANGUAGE C STRICT;

CREATE FUNCTION ifx_fdw_validator(text[], oid) RETURNS void
  AS 'MODULE_PATHNAME'
  LANGUAGE C STRICT;

CREATE FOREIGN DATA WRAPPER informix_fdw
  HANDLER ifx_fdw_handler
  VALIDATOR ifx_fdw_validator;
```

FDW internal

Datum

```
ifx_fdw_handler(PG_FUNCTION_ARGS)
{
    FdwRoutine *fdwRoutine          = makeNode(FdwRoutine);

    /* Required */
    fdwroutine->GetForeignRelSize    = blackholeGetForeignRelSize;    /* S    U D */
    fdwroutine->GetForeignPaths      = blackholeGetForeignPaths;      /* S    U D */
    fdwroutine->GetForeignPlan       = blackholeGetForeignPlan;       /* S    U D */
    fdwroutine->BeginForeignScan     = blackholeBeginForeignScan;     /* S    U D */
    fdwroutine->IterateForeignScan   = blackholeIterateForeignScan;   /* S          */
    fdwroutine->ReScanForeignScan    = blackholeReScanForeignScan;    /* S          */
    fdwroutine->EndForeignScan       = blackholeEndForeignScan;       /* S    U D */

    /* Optional - use NULL if not required */
    ...
    ...

    PG_RETURN_POINTER(fdwRoutine);
}
```

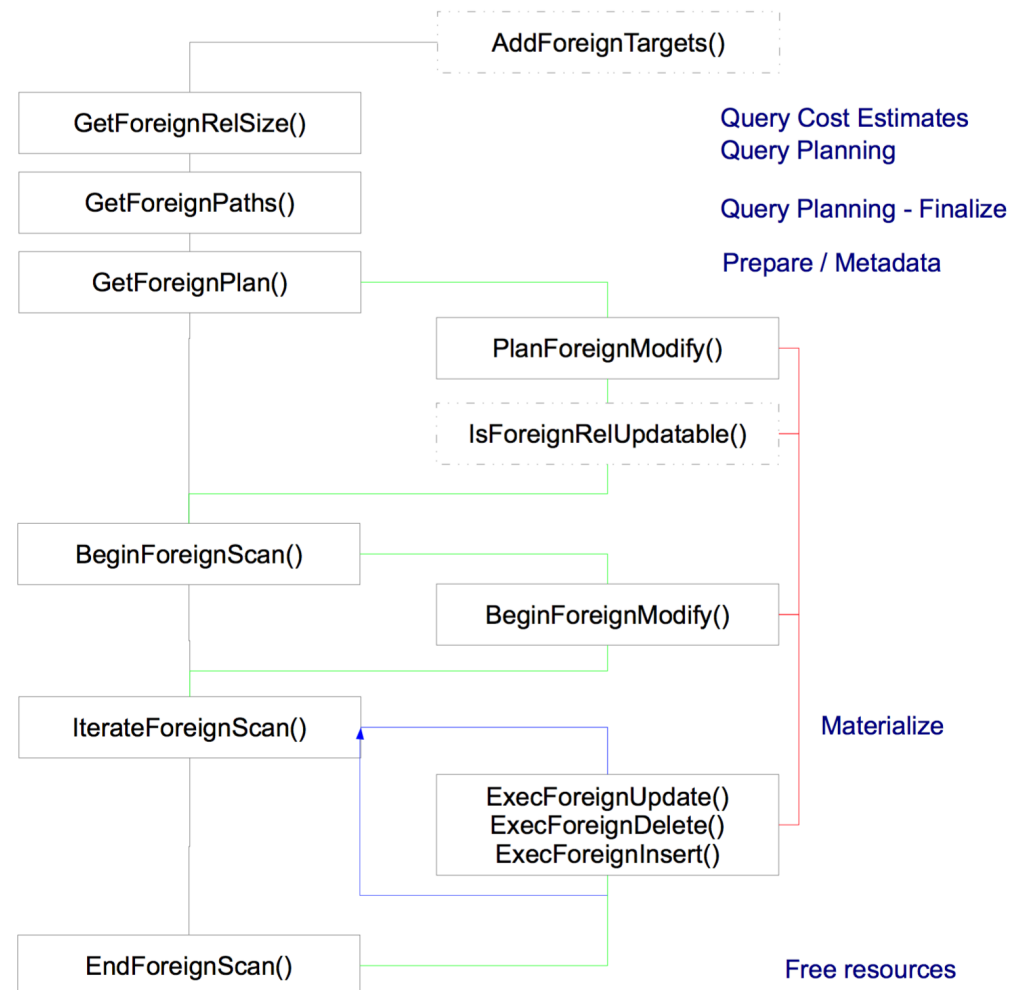
FDW internal

```
/* Optional - use NULL if not required */
#if (PG_VERSION_NUM >= 90300)
    fdwroutine->IsForeignRelUpdatable      = blackholeIsForeignRelUpdatable;
    fdwroutine->AddForeignUpdateTargets     = blackholeAddForeignUpdateTargets; /*      U D */
    fdwroutine->PlanForeignModify           = blackholePlanForeignModify;      /*    I U D */
    fdwroutine->BeginForeignModify          = blackholeBeginForeignModify;      /*    I U D */
    fdwroutine->ExecForeignInsert           = blackholeExecForeignInsert;        /*      I   */
    fdwroutine->ExecForeignUpdate           = blackholeExecForeignUpdate;        /*        U   */
    fdwroutine->ExecForeignDelete           = blackholeExecForeignDelete;        /*          D */
    fdwroutine->EndForeignModify            = blackholeEndForeignModify;        /*    I U D */
#endif

    /* support for EXPLAIN */
    fdwroutine->ExplainForeignScan          = blackholeExplainForeignScan;      /* S      U D EXPLAIN */
    fdwroutine->ExplainForeignModify        = blackholeExplainForeignModify;    /*      I U D EXPLAIN */
    fdwroutine->AnalyzeForeignTable         = blackholeAnalyzeForeignTable;      /* ANALYZE */

#if (PG_VERSION_NUM >= 90500)
    fdwroutine->ImportForeignSchema         = blackholeImportForeignSchema;    /* IMPORT FOREIGN SCHEMA */
    fdwroutine->GetForeignJoinPaths         = blackholeGetForeignJoinPaths;    /* Scanning foreign joins */
    fdwroutine->GetForeignRowMarkType       = blackholeGetForeignRowMarkType;  /* Locking foreign rows */
    fdwroutine->RefetchForeignRow           = blackholeRefetchForeignRow;      /* Locking foreign rows */
#endif
```

FDW Flow



FDW Query Planing

- Setup and Planning a scan or modify action on a foreign datasource
- E.g. establish and cache remote connection
- Initialize required supporting structures for remote access
- Planner info and cost estimates via baserel and root parameters.

`GetForeignRelSize()`

`GetForeignPaths()`

`GetForeignPlan()`

FDW Query Scanning - `BeginForeignScan()`

`void`

```
BeginForeignScan (ForeignScanState *node,  
                  int eflags);
```

- Execute startup callback for the FDW.
- Basically prepares the FDW for executing a scan.
- `ForeignScanState` saves function state values.
- Use `node->fdw_state` to assign your own FDW state structure.
- Must handle EXPLAIN and EXPLAIN ANALYZE by checking `eflags` & `EXEC_FLAG_EXPLAIN_ONLY`

FDW Query Scanning - IterateForeignScan()

```
TupleTableSlot *
```

```
IterateForeignScan (ForeignScanState *node);
```

- Fetches data from the remote source.
- Data conversion
- Materializes a physical or virtual tuple to be returned.
- Needs to return an empty tuple when done.
- Private FDW data located in `node->fdw` state

FDW Query Scanning - `ReScanForeignScan()`

```
void ReScanForeignScan (ForeignScanState *node);
```

- Prepares the FDW to handle a rescan
- Begins the scan from the beginning, e.g when used in scrollable cursors
- Must take care for changed query parameters!

End FDW Query Scanning - **EndForeignScan()**

```
void EndForeignScan (ForeignScanState *node);
```

- **EndForeignScan()** run when **IterateForeignScan** returns no more rows
- Finalizes the remote scan
- Close result sets, handles, connection, free memory, etc...

Memory Management

- PostgreSQL uses `palloc()`
- Memory is allocated in `CurrentMemoryContext`
- Use your own `MemoryContext` where necessary
 - e.g. `IterateForeignScan()`
- Memory allocated in external libraries need special care

More Information

- Writing A Foreign Data Wrapper
 - <http://www.slideshare.net/psoo1978/pg-fdw>
- Official Documentation
 - <http://www.postgresql.org/docs/current/static/index.html>
 - Chapter 35. Extending SQL
 - Chapter 54. Writing A Foreign Data Wrapper
- Wiki page
 - https://wiki.postgresql.org/wiki/Foreign_data_wrappers

file_fdw

- <http://www.postgresql.org/docs/9.5/static/file-fdw.html>
- Can be used to access data files in the server's file system.
- Internally, uses COPY.

```
CREATE EXTENSION file_fdw;  
CREATE SERVER pglog FOREIGN DATA WRAPPER file_fdw;  
CREATE FOREIGN TABLE pglog (  
    log_time timestamp(3) with time zone,  
    user_name text,  
    database_name text,  
    process_id integer,  
    ...  
    application_name text  
) SERVER pglog  
OPTIONS ( filename '/home/josh/9.1/data/pg_log/pglog.csv', format 'csv' );
```


postgres_fdw

- <http://www.postgresql.org/docs/current/static/postgres-fdw.html>
- Similar with dblink module but can give better performance in many cases.

```
CREATE EXTENSION postgres_fdw;

CREATE SERVER foreign_server
    FOREIGN DATA WRAPPER postgres_fdw
    OPTIONS (host '192.83.123.89', port '5432', dbname 'foreign_db');

CREATE USER MAPPING FOR local_user
    SERVER foreign_server
    OPTIONS (user 'foreign_user', password 'password');

CREATE FOREIGN TABLE foreign_table (
    id integer NOT NULL,
    data text
)
    SERVER foreign_server
    OPTIONS (schema_name 'some_schema', table_name 'some_table');
```

oracle_fdw

- https://github.com/laurenz/oracle_fdw

```
pgdb=# CREATE EXTENSION oracle_fdw;

pgdb=# CREATE SERVER oradb FOREIGN DATA WRAPPER oracle_fdw
        OPTIONS (dbserver '//dbserver.mydomain.com/ORADB');

pgdb=# GRANT USAGE ON FOREIGN SERVER oradb TO pguser;

pgdb=> CREATE USER MAPPING FOR pguser SERVER oradb
        OPTIONS (user 'orauser', password 'orapwd');

pgdb=> CREATE FOREIGN TABLE oratab (
        id          integer          OPTIONS (key 'true') NOT NULL,
        text        character varying(30),
        floating    double precision NOT NULL
    ) SERVER oradb OPTIONS (schema 'ORAUSER', table 'ORATAB');
```

Lab